


Five arms: only the split position and the decoder width change



Held constant across all five: encoder, block structure, L1 loss with 1.0/1.0 weights, Adam 1e-3 cosine to 1e-5, batch 16, ≤200 epochs with early stopping, no augmentation, one dataset-wide normalisation, and the same 58 seeds on the same data.

Decoder block height is drawn proportional to width multiplier, so the wide arms are visibly wider. D against B_wide is the primary contrast: 10,811,298 against 10,743,738 parameters, a 0.62% difference, identical supervision and identical data — so the only thing left to explain a difference would be whether the decoder is split. Arm A (one output, historical baseline) is dropped here: it has no heart head and cannot be scored on the primary endpoint.