



The heart reconstruction is the endpoint, and on the lung's scale you can barely see it


(a) input, 40 dB
lung + heart, registered level



(b) lung reconstruction
target, shared scale




(c) heart reconstruction
SAME scale as (a),(b)




*almost invisible beside
the lung — the premise*

(d) heart reconstruction
its OWN scale




*peak 5.8 vs the lung's
17.0 — 34% of it*

(e) input, 20 dB
where H3 appears



(f) (e) — (a)
peak 5.4% of (a)



reconstructed conductivity change (GREIT units) — shared scale

One test-split geometry (index 1945), the median heart radius of the test split. Every panel is an EIDORS GREIT reconstruction: `gen_dataset.m` runs three forward and three inverse solves per sample (mixed, lung only, heart only).
(c) and (d) are the SAME reconstruction on two scales — that gap is what the architecture claims to address, and heart DSC is the registered primary endpoint of study 2.
(f) GREIT is a regularised inverse and suppresses measurement noise before image space: 40 dB to 20 dB changes the input by 5.4% of peak yet costs 8.8 DSC points (arm B_wide, 58 seeds). Bars are model units; 64 px = 4 units.